

Linking

CS 47: Introduction to Computer Systems
April 9 and 14, 2014

Original lecture by:

Randy Bryant and Dave O'Hallaron

Modified for SJSU by Thomas D. Howell

CS 47 Spring 2014

1

Example C Program

main.c

```
int buf[2] = {1, 2};

int main()
{
    swap();
    return 0;
}
```

swap.c

```
extern int buf[];

int *bufp0 = &buf[0];
static int *bufp1;

void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

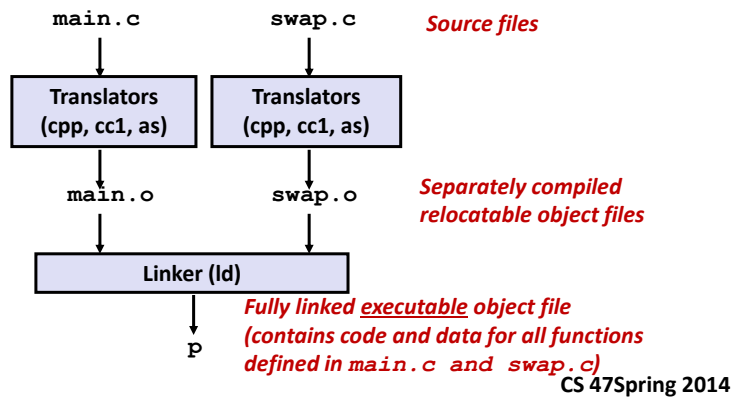
CS 47Spring 2014

2

Static Linking

■ Programs are translated and linked using a *compiler driver*:

- `unix> gcc -O2 -g -o p main.c swap.c`
- `unix> ./p`



3

Why Linkers?

■ Reason 1: Modularity

- Program can be written as a collection of smaller source files, rather than one monolithic mass.
- Can build libraries of common functions (more on this later)
 - e.g., Math library, standard C library

4

Why Linkers? (cont)

■ Reason 2: Efficiency

- Time: Separate compilation
 - Change one source file, compile, and then relink.
 - No need to recompile other source files.
- Space: Libraries
 - Common functions can be aggregated into a single file...
 - Yet executable files and running memory images contain only code for the functions they actually use.

CS 47Spring 2014

5

What Do Linkers Do?

■ Step 1. Symbol resolution

- Programs define and reference *symbols* (variables and functions):
 - `void swap() {...}` */* define symbol swap */*
 - `swap();` */* reference symbol a */*
 - `int *xp = &x;` */* define symbol xp, reference x */*
- Symbol definitions are stored (by compiler) in *symbol table*.
 - Symbol table is an array of structs
 - Each entry includes name, size, and location of symbol.
- Linker associates each symbol reference with exactly one symbol definition.

CS 47Spring 2014

6

What Do Linkers Do? (cont)

■ Step 2. Relocation

- Merges separate code and data sections into single sections
- Relocates symbols from their relative locations in the `.o` files to their final absolute memory locations in the executable.
- Updates all references to these symbols to reflect their new positions.

CS 47Spring 2014

7

Three Kinds of Object Files (Modules)

■ Relocatable object file (`.o` file)

- Contains code and data in a form that can be combined with other relocatable object files to form executable object file.
 - Each `.o` file is produced from exactly one source (`.c`) file

■ Executable object file (`a.out` [`a.exe`] file)

- Contains code and data in a form that can be copied directly into memory and then executed.

■ Shared object file (`.so` file)

- Special type of relocatable object file that can be loaded into memory and linked dynamically, at either load time or run-time.
- Called *Dynamic Link Libraries* (DLLs) by Windows

CS 47Spring 2014

8

Executable and Linkable Format (ELF)

- **Standard binary format for object files**
- **Originally proposed by AT&T System V Unix**
 - Later adopted by BSD Unix variants and Linux
- **One unified format for**
 - Relocatable object files (`.o`),
 - Executable object files (`a.out`)
 - Shared object files (`.so`)
- **Generic name: ELF binaries**
- **Windows version is called COFF (Common Object File Format). Similar but incompatible.**

CS 47Spring 2014

9

ELF Object File Format

- **Elf header**
 - Word size, byte ordering, file type (`.o`, `exec`, `.so`), machine type, etc.
- **Segment header table**
 - Page size, virtual addresses memory segments (sections), segment sizes.
- **.text section**
 - Code
- **.rodata section**
 - Read only data: jump tables, ...
- **.data section**
 - Initialized global variables
- **.bss section**
 - Uninitialized global variables
 - "Block Storage Start"
 - "Better Save Space"
 - Has section header but occupies no space

0	ELF header
	Segment header table (required for executables)
	.text section
	.rodata section
	.data section
	.bss section
	.symtab section
	.rel.text section
	.rel.data section
	.debug section
	Section header table

CS 47Spring 2014

10

ELF Object File Format (cont.)

- **.symtab section**
 - Symbol table
 - Procedure and static variable names
 - Section names and locations
- **.rel.text section**
 - Relocation info for **.text** section
 - Addresses of instructions that will need to be modified in the executable
 - Instructions for modifying.
- **.rel.data section**
 - Relocation info for **.data** section
 - Addresses of pointer data that will need to be modified in the merged executable
- **.debug section**
 - Info for symbolic debugging (**gcc -g**)
- **Section header table**
 - Offsets and sizes of each section

ELF header
Segment header table (required for executables)
.text section
.rodata section
.data section
.bss section
.symtab section
.rel.txt section
.rel.data section
.debug section
Section header table

CS 47Spring 2014

11

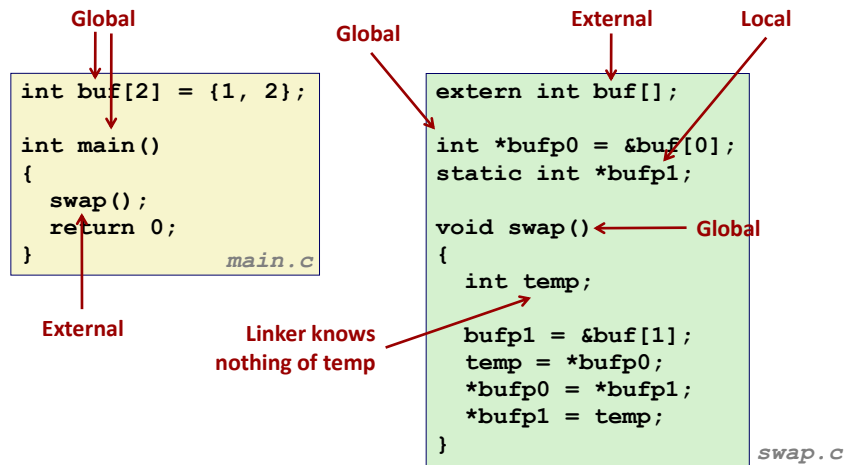
Linker Symbols

- **Global symbols**
 - Symbols defined by module *m* that can be referenced by other modules.
 - E.g.: non-**static** C functions and non-**static** global variables.
- **External symbols**
 - Global symbols that are referenced by module *m* but defined by some other module. (e.g. sleep)
- **Local symbols**
 - Symbols that are defined and referenced exclusively by module *m*.
 - E.g.: C functions and variables defined with the **static** attribute.
 - **Local linker symbols are not local program variables**

CS 47Spring 2014

12

Resolving Symbols



CS 47Spring 2014

13

Practice Problem 7.1

Using program from previous slide:

Symbol	In Swap.o symtab?	Symbol type	Module where defined	Section
buf		External		
bufp0		Global		.data
bufp1	Yes	Local	swap.o	.bss
swap				.text
temp	No			

Note: `bufp1` is *not* declared "static" in Fig. 7.1(b), p. 656


CS 47Spring 2014

14

Strong and Weak Symbols

- Program symbols are either strong or weak

- **Strong**: procedures and initialized globals
- **Weak**: uninitialized globals



CS 47Spring 2014

15

Linker's Symbol Rules

- Rule 1: Multiple strong symbols are not allowed

- Each item can be defined only once
- Otherwise: Linker error

- Rule 2: Given a strong symbol and multiple weak symbol, choose the strong symbol

- References to the weak symbol resolve to the strong symbol

- Rule 3: If there are multiple weak symbols, pick an arbitrary one

- Can override this with `gcc -fno-common`

CS 47Spring 2014

16

Linker Puzzles

```
int x;
p1() {}
```

```
p1() {}
```

Link time error: two strong symbols (**p1**)

```
int x;
p1() {}
```

```
int x;
p2() {}
```

References to **x** will refer to the same uninitialized int. Is this what you really want?

```
int x;
int y;
p1() {}
```

```
double x;
p2() {}
```

Writes to **x** in **p2** might overwrite **y**!
Evil!

```
int x=7;
int y=5;
p1() {}
```

```
double x;
p2() {}
```

Writes to **x** in **p2** will overwrite **y**!
Nasty!

```
int x=7;
p1() {}
```

```
int x;
p2() {}
```

References to **x** will refer to the same initialized variable.

Nightmare scenario: two identical weak structs, compiled by different compilers with different alignment rules. CS 47Spring 2014

17

Role of .h Files

c1.c

```
#include "global.h"

int f() {
    return g+1;
}
```

c2.c

```
#include <stdio.h>
#include "global.h"

int main() {
    if (!init)
        g = 37;
    int t = f();
    printf("Calling f yields %d\n", t);
    return 0;
}
```

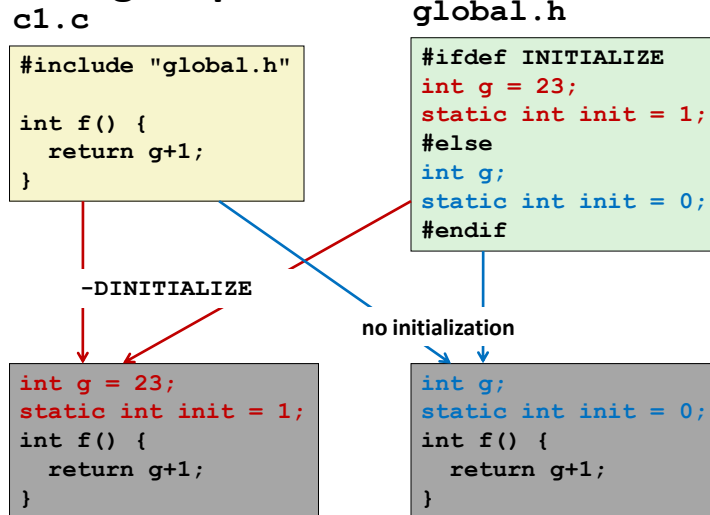
global.h

```
#ifndef INITIALIZE
int g = 23;
static int init = 1;
#else
int g;
static int init = 0;
#endif
```

CS 47Spring 2014

18

Running Preprocessor

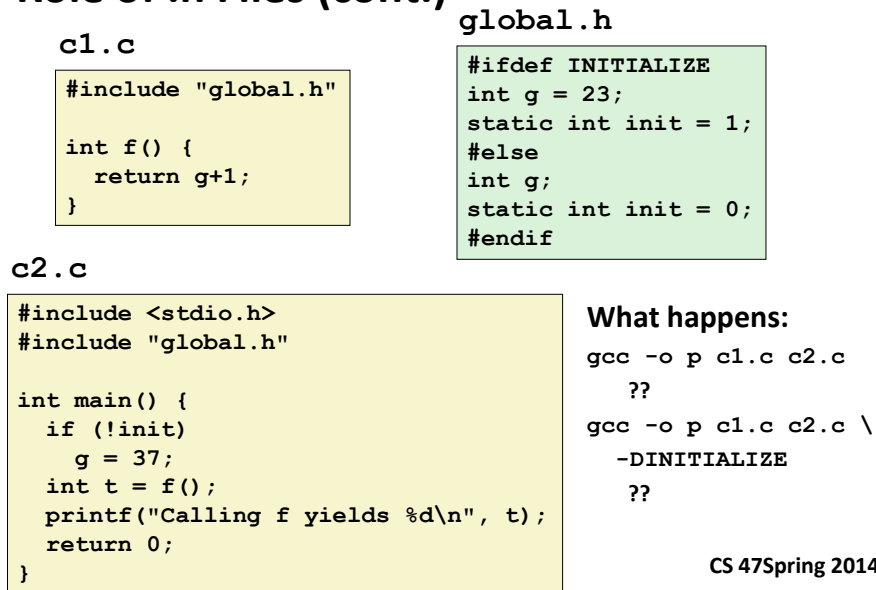


CS 47Spring 2014

#include causes C preprocessor to insert file verbatim

19

Role of .h Files (cont.)



CS 47Spring 2014

20

Global Variables

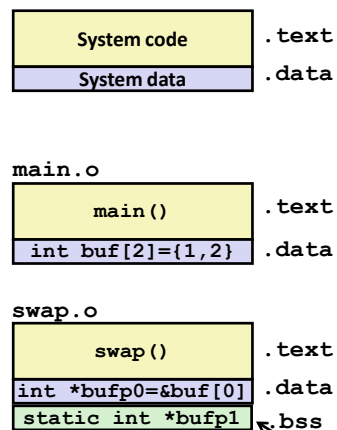
- Avoid if you can
- Otherwise
 - Use `static` if you can
 - Initialize if you define a global variable
 - Use `extern` if you use external global variable

CS 47Spring 2014

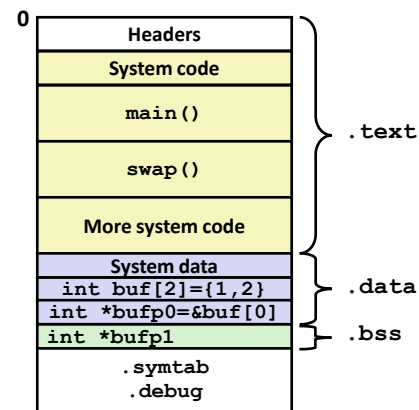
21

Relocating Code and Data

Relocatable Object Files



Executable Object File



Even though private to swap, requires allocation in .bss

CS 47Spring 2014

22

Relocation Info (main)

main.c

```
int buf[2] =
{1,2};

int main()
{
    swap();
    return 0;
}
```

main.o

```
00000000 <_main>:
0: 55                push    %ebp
1: b8 10 00 00 00    mov     $0x10,%eax
6: 89 e5             mov     %esp,%ebp
8: 83 ec 08          sub     $0x8,%esp
b: 83 e4 f0          and     $0xffffffff0,%esp
e: e8 00 00 00 00    call   13 <_main+0x13>
                        f: DISP32    __alloca
13: e8 00 00 00 00    call   18 <_main+0x18>
                        14: DISP32    __main
18: e8 00 00 00 00    call   1d <_main+0x1d>
                        19: DISP32    __swap
1d: c9                leave   %ebp
1e: 31 c0             xor     %eax,%eax
20: c3                ret
```

Disassembly of section .data:

```
00000000 <_buf>:
0: 01 00 00 00
4: 02 00 00 00
```

Source: objdump -r -D

CS 47Spring 2014

23

Relocation Info (swap, .text)

swap.c

```
extern int buf[];

int
*bufp0 = &buf[0];

static int *bufp1;

void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

swap.o

Disassembly of section .text:

```
00000000 <_swap>:
0: 55                push    %ebp
1: b8 04 00 00 00    mov     $0x4,%eax    # &buf[1]
2: dir32    _buf
6: 8b 15 04 00 00 00    mov     0x4,%edx    # buf[1]
8: dir32    _buf
c: a3 00 00 00 00    mov     %eax,0x0    # bufp1
d: dir32    .bss
11: a1 00 00 00 00    mov     0x0,%eax    # bufp0
12: dir32    .data
16: 89 e5             mov     %esp,%ebp
18: 8b 08             mov     (%eax),%ecx  # temp
1a: 89 10             mov     %edx,(%eax)  # *bufp0
1c: 5d                pop     %ebp
1d: 89 0d 04 00 00 00    mov     %ecx,0x4    # *bufp1
1f: dir32    _buf
23: c3                ret
```

24

Relocation Info (swap, .data)

swap.c

```
extern int buf[];

int *bufp0 =
    &buf[0];
static int *bufp1;

void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

Disassembly of section .data:

```
00000000 <_bufp0>:
    0:  00 00 00 00
    0:  dir32      _buf
    ...
```

Disassembly of section .bss:

```
00000000 <_bufp1>:
```

CS 47Spring 2014

25

Executable Before/After Relocation (.text)

00000000 <main>:

```
    . . .
18:  e8 00 00 00 00      call    1d <_main+0x1d>
    . . .                19: DISP32      _swap
1d:  c9                  leave
    . . .
```

0x40105d + 0x13
= 0x401070

00401040 <_main>:

```
401040:  55                  push    %ebp
401041:  b8 10 00 00 00      mov     $0x10,%eax
401046:  89 e5               mov     %esp,%ebp
401048:  83 ec 08             sub     $0x8,%esp
40104b:  83 e4 f0             and     $0xffffffff0,%esp
40104e:  e8 4d 00 00 00      call   4010a0 <__chkstk>
401053:  e8 a8 00 00 00      call   401100 <__main>
401058:  e8 13 00 00 00      call   401070 <_swap>
40105d:  c9                  leave
40105e:  31 c0               xor     %eax,%eax
401060:  c3                  ret
```

CS 47Spring 2014

26

```

1:  b8 04 00 00 00      mov    $0x4,%eax
      2:  dir32      _buf
6:  8b 15 04 00 00 00    mov    0x4,%edx
      8:  dir32      _buf
c:  a3 00 00 00 00      mov    %eax,0x0
      d:  dir32      .bss
11: a1 00 00 00 00      mov    0x0,%eax
      12: dir32      .data
...
1d: 89 0d 04 00 00 00    mov    %ecx,0x4
      1f: dir32      _buf
23: c3                  ret

```

```

00401070 <_swap>:
401070: 55                  push   %ebp
401071: b8 04 20 40 00      mov    $0x402004,%eax
401076: 8b 15 04 20 40 00    mov    0x402004,%edx
40107c: a3 08 30 40 00      mov    %eax,0x403008
401081: a1 08 20 40 00      mov    0x402008,%eax
401086: 89 e5               mov    %esp,%ebp
401088: 8b 08               mov    (%eax),%ecx
40108a: 89 10               mov    %edx,(%eax)
40108c: 5d                  pop    %ebp
40108d: 89 0d 04 20 40 00    mov    %ecx,0x402004
401093: c3                  ret

```

27

Executable After Relocation (.data)

```

Disassembly of section .data:

402000 <buf>:
402000: 01 00 00 00 02 00 00 00
402008 <bufp0>:
402008: 00 20 40 00

```

Packaging Commonly Used Functions

- **How to package functions commonly used by programmers?**
 - Math, I/O, memory management, string manipulation, etc.
- **Awkward, given the linker framework so far:**
 - **Option 1:** Put all functions into a single source file
 - Programmers link big object file into their programs
 - Space and time inefficient
 - **Option 2:** Put each function in a separate source file
 - Programmers explicitly link appropriate binaries into their programs
 - More efficient, but burdensome on the programmer

CS 47Spring 2014

29

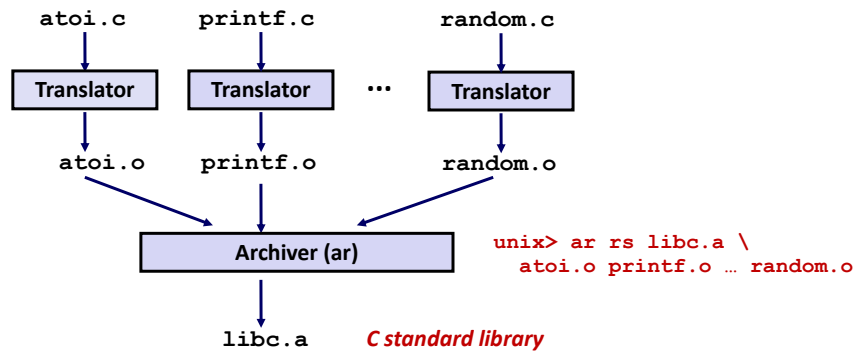
Solution: Static Libraries

- **Static libraries (.a archive files)**
 - Concatenate related relocatable object files into a single file with an index (called an *archive*).
 - Enhance linker so that it tries to resolve unresolved external references by looking for the symbols in one or more archives.
 - If an archive member file resolves reference, link it into the executable.

CS 47Spring 2014

30

Creating Static Libraries



- Archiver allows incremental updates
 - Recompile function that changes and replace .o file in archive.
- CS 47Spring 2014

31

Commonly Used Libraries

libc.a (the C standard library)

- 8 MB archive of 1392 object files.
- I/O, memory allocation, signal handling, string handling, data and time, random numbers, integer math

libm.a (the C math library)

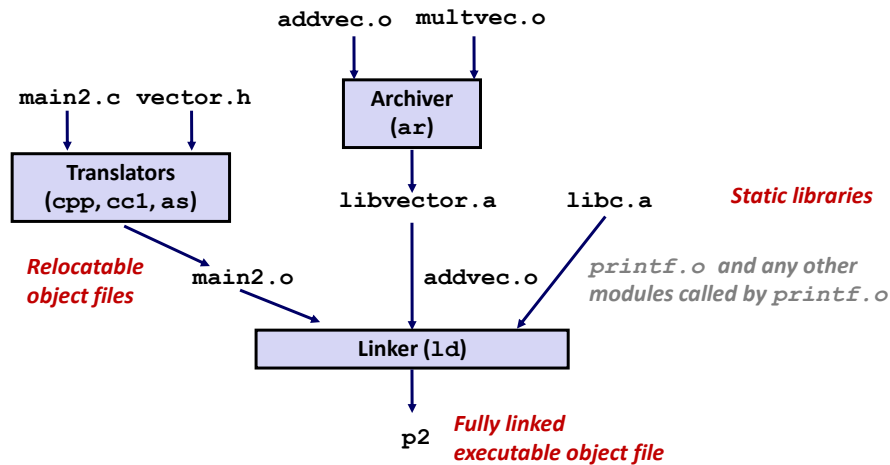
- 1 MB archive of 401 object files.
- floating point math (sin, cos, tan, log, exp, sqrt, ...)

```
% ar -t /usr/lib/libc.a | sort
...
fork.o
...
fprintf.o
fpu_control.o
fputc.o
freopen.o
fscanf.o
fseek.o
fstab.o
...
```

```
% ar -t /usr/lib/libm.a | sort
...
e_acos.o
e_acosf.o
e_acosh.o
e_acoshf.o
e_acoshl.o
e_acosl.o
e_asin.o
e_asinf.o
e_asinl.o
...
```

32

Linking with Static Libraries



CS 47Spring 2014

33

Using Static Libraries

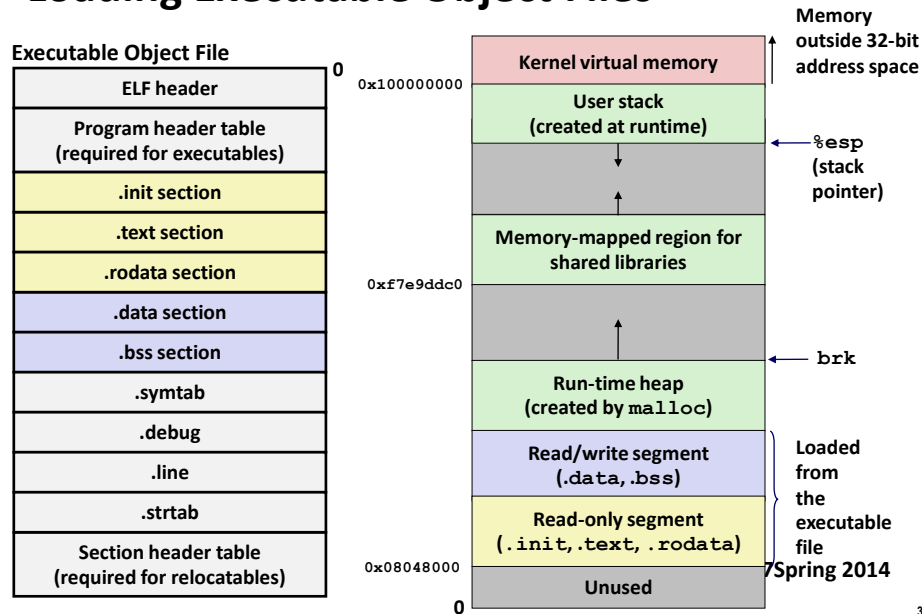
- **Linker's algorithm for resolving external references:**
 - Scan `.o` files and `.a` files in the command line order.
 - During the scan, keep a list of the current unresolved references.
 - As each new `.o` or `.a` file, *obj*, is encountered, try to resolve each unresolved reference in the list against the symbols defined in *obj*.
 - If any entries in the unresolved list at end of scan, then error.
- **Problem:**
 - Command line order matters!
 - Moral: put libraries at the end of the command line.

```
unix> gcc -L. libtest.o -lmine
unix> gcc -L. -lmine libtest.o
libtest.o: In function `main':
libtest.o(.text+0x4): undefined reference to `libfun'
```

CS 47Spring 2014

34

Loading Executable Object Files



35

Shared Libraries

- **Static libraries have the following disadvantages:**
 - Duplication in the stored executables (every function need std libc)
 - Duplication in the running executables
 - Minor bug fixes of system libraries require each application to explicitly relink
- **Modern solution: Shared Libraries**
 - Object files that contain code and data that are loaded and linked into an application *dynamically*, at either *load-time* or *run-time*
 - Also called: dynamic link libraries, DLLs, .so files

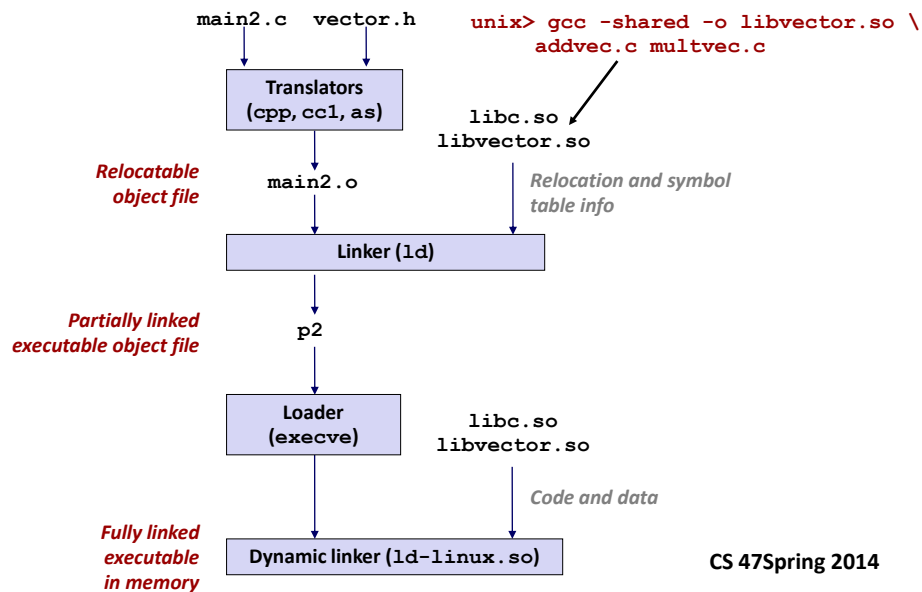
Shared Libraries (cont.)

- **Dynamic linking can occur when executable is first loaded and run (load-time linking).**
 - Common case for Linux, handled automatically by the dynamic linker (`ld-linux.so`).
 - Standard C library (`libc.so`) usually dynamically linked.
- **Dynamic linking can also occur after program has begun (run-time linking).**
 - In Linux, this is done by calls to the `dlopen()` interface.
 - Distributing software.
 - High-performance web servers.
 - Runtime library interpositioning.
- **Shared library routines can be shared by multiple processes.**
 - More on this when we learn about virtual memory

CS 47Spring 2014

37

Dynamic Linking at Load-time



CS 47Spring 2014

38

Dynamic Linking at Run-time

```
#include <stdio.h>
#include <dlfcn.h>

int x[2] = {1, 2};
int y[2] = {3, 4};
int z[2];

int main()
{
    void *handle;
    void (*addvec)(int *, int *, int *, int);
    char *error;

    /* dynamically load the shared lib that contains addvec() */
    handle = dlopen("./libvector.so", RTLD_LAZY);
    if (!handle) {
        fprintf(stderr, "%s\n", dlerror());
        exit(1);
    }
}
```

CS 47Spring 2014

39

Dynamic Linking at Run-time

```
...

/* get a pointer to the addvec() function we just loaded */
addvec = dlsym(handle, "addvec");
if ((error = dlerror()) != NULL) {
    fprintf(stderr, "%s\n", error);
    exit(1);
}

/* Now we can call addvec() just like any other function */
addvec(x, y, z, 2);
printf("z = [%d %d]\n", z[0], z[1]);

/* unload the shared library */
if (dlclose(handle) < 0) {
    fprintf(stderr, "%s\n", dlerror());
    exit(1);
}
return 0;
}
```

CS 47Spring 2014

40